

Topics on High-Dimensional Data Analytics

Regularization Applications

Kamran Paynabar, Ph.D.

Associate Professor
School of Industrial and Systems Engineering

Compressive Sensing I



Learning Objectives


- Explain principles of compressive sensing
- Use L0 and L1 regularization to formulate a CS problem.
- Apply optimization algorithms to solve a CS problem.



Sparse Information and Regularization

- High-dimensional data usually have a low dimensional structure.
- Important information of HD data is embedded in a few dimensions (sparse) and the rest are non-informative and noise.


Single-channel signals




Multi-channel signals



Images




Point Cloud



- Regularization helps identify the sparsely informative features and remove the noise.


Sparse Representation of Image

- Image Compression-JPEG



compression:
JPEG, JPEG2000, MPEG, ...

- Wavelet




$$x = \sum_{i=1}^N \alpha_i \psi_i$$

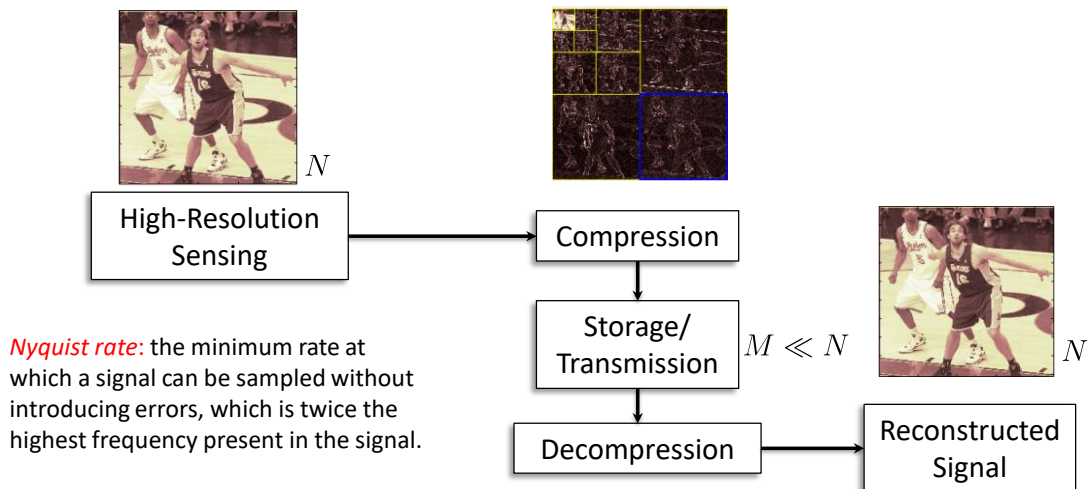
$\alpha_i = \langle x, \psi_i \rangle$

Image Compression

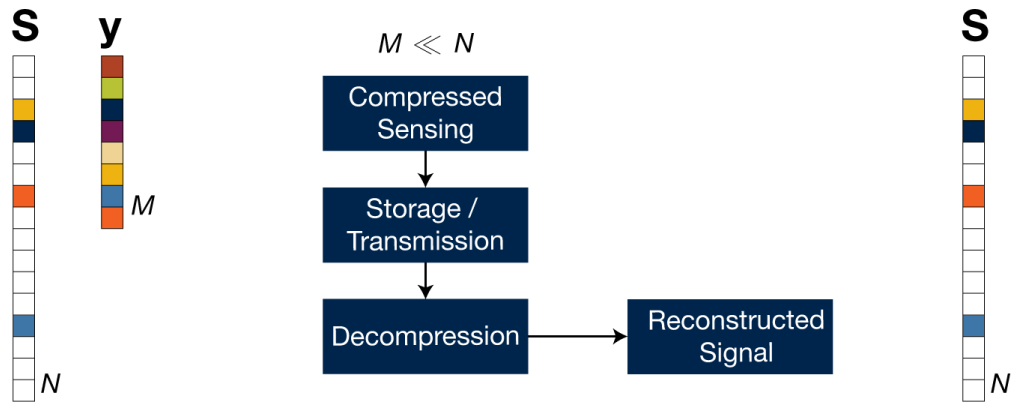
- Keep the 1% **largest** coefficients and set the rest to zero.



Conventional Sampling & Compression



Compressed Sampling



The original or transformed signal is sparse (a few non-zero elements).

A Simple Example

- Consider a 1-support signal (i.e., all zeros but one).
- Using Hamming Matrix the signal can be measured in a compressed way

$$\begin{bmatrix} 5 \\ 5 \\ 0 \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 & 1 \end{bmatrix} \times \begin{bmatrix} 0 \\ 5 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$


- Only $\log_2 n$ measurement is needed to recover the signal

Sensing Matrix


- Redefine sensing approach by using sensing matrix
- Sensing Matrix
 - Each measurement is a vector ϕ_i of dimension n
 - Given ϕ_i and signal, measurement $y_i = \phi_i^T s$
 - Cost is proportional to the number of measurements

1	1	1	1	1	1	1	1
1	1	1	1	0	0	0	0
1	1	0	0	1	1	0	0
1	0	1	0	1	0	1	0

Hamming Matrix



Random Matrix



1	0	0	...	0
0	1	0	...	0
0	0	1	...	0
⋮	⋮	⋮	⋱	⋮
0	0	0	...	1

Identity Matrix



Traditional vs Compressive Sensing Matrix

$$\begin{bmatrix} Y_1 \\ \vdots \\ Y_n \end{bmatrix} = \begin{bmatrix} \text{Diagonal Matrix} \end{bmatrix} \begin{bmatrix} f_1^* \\ f_2^* \\ \vdots \\ f_n^* \end{bmatrix}$$

Sample Uniformly at a high rate (Nyquist Rate)

$$\begin{bmatrix} Y_1 \\ \vdots \\ Y_k \end{bmatrix} = \begin{bmatrix} \text{Random Matrix} \end{bmatrix} \begin{bmatrix} f_1^* \\ f_2^* \\ \vdots \\ f_n^* \end{bmatrix}$$

Sample randomly, and Take very few samples

Examples of Random Sensing Matrices



Gaussian random matrix



Rademacher matrix
(i.i.d. of +1 and -1 entries)



Fourier random matrix
(rows are randomly selected DFT vectors)

Topics on High-Dimensional Data Analytics

Regularization Applications

Kamran Paynabar, Ph.D.


Associate Professor
School of Industrial and Systems Engineering

Compressive Sensing II




Learning Objectives

- Explain principles of compressive sensing
- Use L0 and L1 regularization to formulate a CS problem.
- Apply optimization algorithms to solve a CS problem.



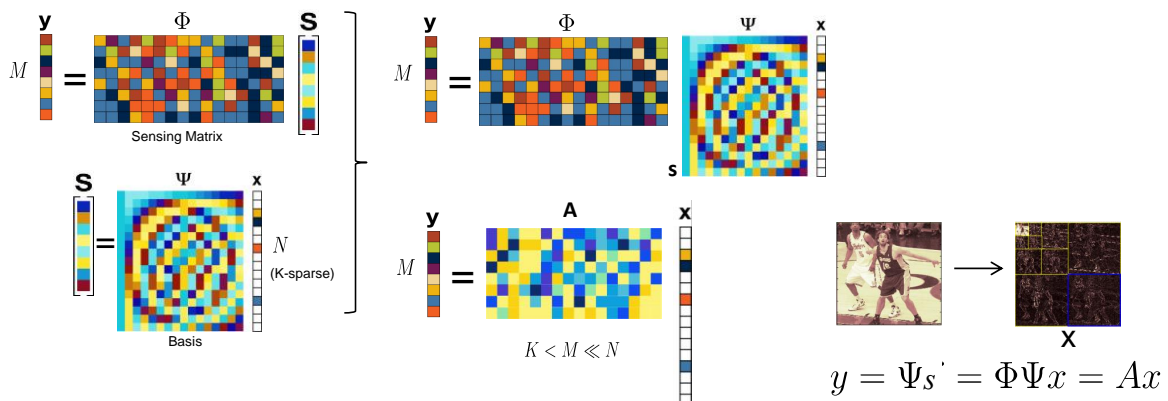
Compressive Sensing Framework

If the signal is sparse in original domain, then we use a fat sensing matrix for sampling.



Compressive Sensing Framework ^s

If the signal is NOT sparse in original domain, we should transform it into a domain with sparse representation.



Signal Reconstruction from Measurements

- We should solve an underdetermined linear system $y = Ax$

$$y = Ax$$

M $K < M \ll N$

- Search for the sparsest signal that agrees with the measurements


$$\text{minimize } \|x\|_0 \quad \text{subject to } y = Ax.$$

$$\|x\|_0 \doteq \#\{i \mid x_i \neq 0\}$$

Non-convex and Interactable

Relaxation of L_0 by L_2


$$\|x\|_p = (\sum_i |x_i|^p)^{1/p}$$



Problem is convex when $p \geq 1$

L_2 Norm: fast but **inaccurate** results

$$\hat{x} = \arg \min_{y=Ax} \|x\|_2$$




x



$$\hat{x} = (A^T A)^{-1} A^T y$$

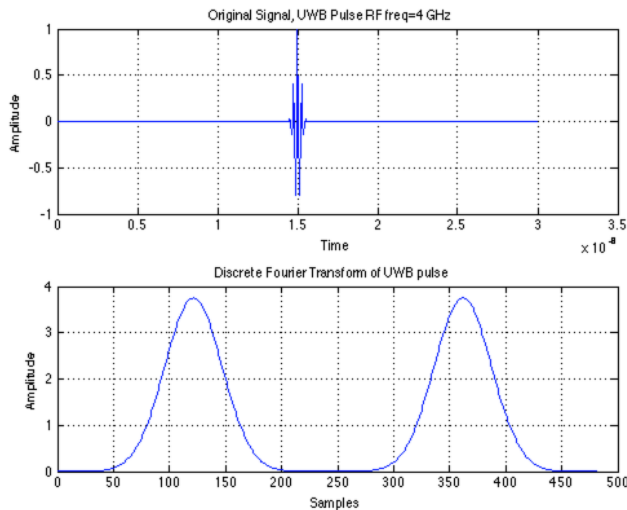
Relaxation of L_0 by L_1

minimize $\|x\|_1$ subject to $Ax = y$. **Efficiently solvable**



Norm Used	Property	Formulation
L2	Fast, Wrong	$\hat{x} = \arg \min_{y=Ax} \ x\ _2$
L0	Slow, Correct	$\hat{x} = \arg \min_{y=Ax} \ x\ _0$
L1	Fast, Mild Oversampling	$\hat{x} = \arg \min_{y=Ax} \ x\ _1$

Example 1: Sparse in Time Domain

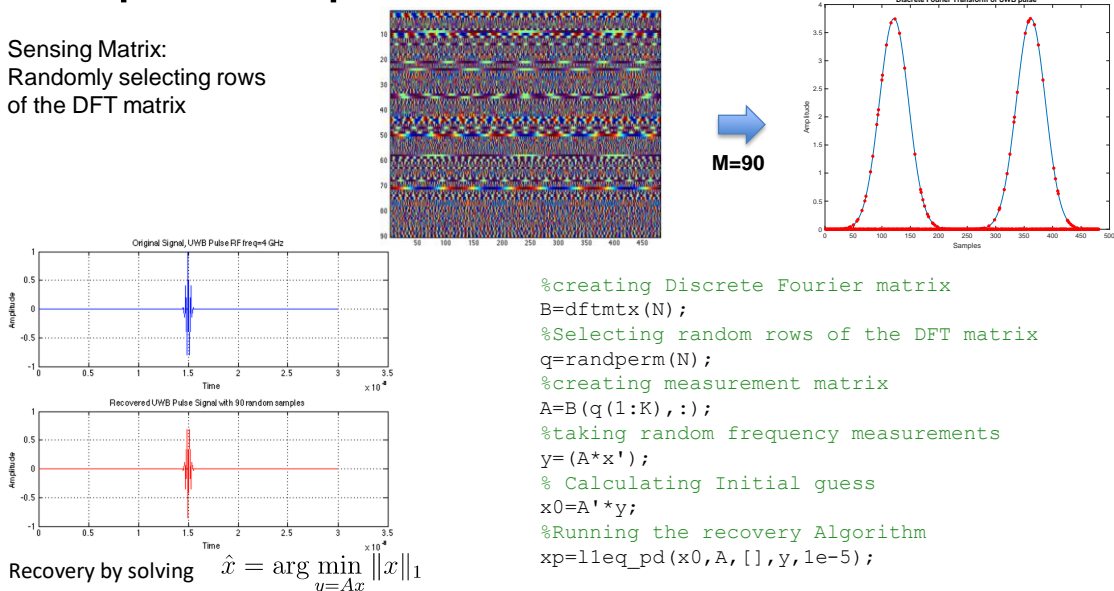


Generate Signal by
`x=pulstran(t,15e-9,'gauspuls',f,0.5);`

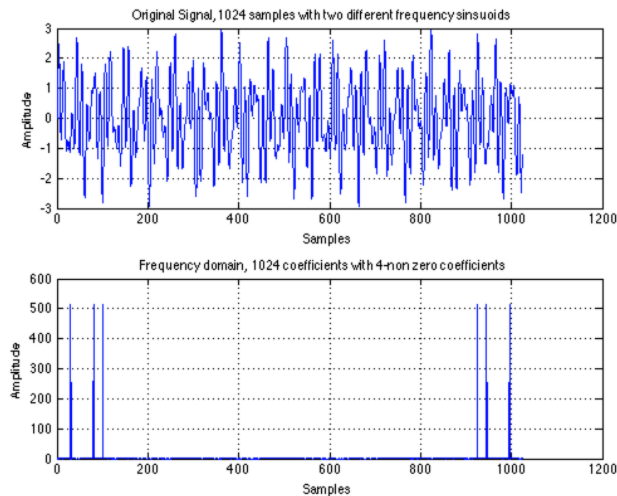
FFT analysis
`xf=fft(x);`

Example 1: Sparse in Time Domain

Sensing Matrix:
 Randomly selecting rows
 of the DFT matrix



Example 2: Sparse in Frequency Domain




Generate Signal by

```
x=(sin(2*pi*(k1/N)*n)+sin(2*pi*(k2/N)*n)+sin(2*pi*(k3/N)*n));
```

FFT analysis

```
xf=fft(x);
```

Example 2: Sparse in Frequency Domain





Φ Sensing Matrix:

Random rows of
identity matrix
with
M=650


Ψ Fourier Basis $y = \Psi s' = \Phi \Psi x = Ax$

```
ID = eye(N); q=randperm(N);
Phi=IN(:,q(1:K))';
Psi = dftmtx(N);
%taking random time measurements
y=(Phi*x);
%Calculating Initial guess
x0= Psi'*(Phi'*y);
%Running the recovery Algorithm
xp=l1eq_pd(x0,Phi*Psi,[],y,1e-7);
%recovered signal in time domain
xprec=real(-inv(Psi)*xp);
```



Recovery by solving $\hat{x} = \arg \min_{y=Ax} \|x\|_1$

Example 2: Sparse in Frequency Domain



Exact
Recovery

```
%recovered signal in time domain
xprec=real(-inv(Psi)*xp)
```

Example 3: CS Applications for Images

Take $M=48640$ incoherent measurements


$f_a = \Psi x$ wavelet coefficients (perfectly sparse)

Solve $\min \|x\|_1 \quad s.t. \quad \Phi \Psi x = y$



original Image

256*256=65536 pixels



Transferred Image

Sparse Wavelet Approximation

Wavelet_OMP_SD_main.m

Example 3: CS Applications for Images



Recover from ~49k Random Measurement

Wavelet_OMP_SD_main.m

Incoherent Sampling

- Coherence between sensing basis Φ and representation basis Ψ


$$\mu(\Phi, \Psi) = \sqrt{n} \cdot \max_{1 \leq k, j \leq n} |\langle \varphi_k, \psi_j \rangle|.$$

- If these matrices contain correlated elements, the coherence is large. Otherwise, it is small.
- It can be shown that

$$\mu(\Phi, \Psi) \in [1, \sqrt{n}]$$


- Compressive Sampling is mainly concerned with low coherence pairs.

Incoherent Pairs: Time-Frequency



Identity Matrix


$$\mu(\Phi, \Psi) = 1$$




Fourier Series

Maximum Incoherence

Incoherent Pairs: Wavelet and Noiselet



Haar basis



Noiselet basis

$$\mu(\Phi, \Psi) = \sqrt{2}$$

Guaranteed Recovery

- Theorem: Fix f and suppose that the coefficient sequence x of f in the basis Ψ is S -Sparse. Select m measurements in Φ domain uniformly at random, then if

$$m \geq C \cdot \mu^2(\Phi, \Psi) \cdot S \cdot \log n$$

holds for some positive constant C , the solution of L_1 formulation is exact with overwhelming probability.

- Remarks
 - Role of coherence is important, smaller the coherence, the fewer samples are needed.
 - If incoherence is close to 1, only $(S \log n)$ samples are needed instead of n .

(Candès and Tao, '04; Donoho, '04):

Signal Recovery from Noisy Measurements and RIP

$$\min \|x\|_{l_1} \quad \text{subject to} \quad \|Ax - y\|_{l_2} \leq \epsilon$$


- If **restricted isometry property (RIP)** holds for the sensing matrix, the signal can be recovered from noisy measurement.
- For each integer $S=1,2,\dots$, define the isometry constant δ_S of a matrix A as the smallest number such that

$$(1 - \delta_S) \|x\|_{\ell_2}^2 \leq \|Ax\|_{\ell_2}^2 \leq (1 + \delta_S) \|x\|_{\ell_2}^2$$

holds for all S -sparse vectors x

- What does it mean?
 - Matrix A Preserve the Euclidean length of S -sparse signals
 - Any S columns taken from A are in fact nearly orthogonal

Example: Noisy Signal Recovery



Recovered by solving $\min \|x\|_{l_1}$ subject to $\|Ax - y\|_{l_2} \leq \epsilon$
 N=512 and M=150 1dnoise_example.m

Example: Noisy Image Recovery




Image Size: 256*256=65K Recovered from 20K measurements

Topics on High-Dimensional Data Analytics

Regularization Applications

Kamran Paynabar, Ph.D.

Associate Professor
School of Industrial and Systems Engineering

Matrix Completion



Learning Objectives

- Explain principles of matrix completion for missing data imputation.
- Use nuclear regularization to formulate a matrix completion problem.
- Apply optimization algorithms to solve a the matrix completion problem.



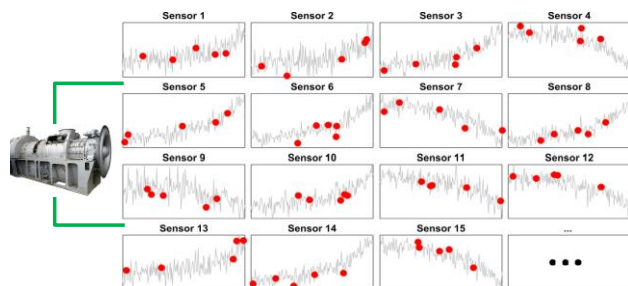
Motivation: Netflix Prize

- Training Data: ~480K users, ~18K movies, 100 M ratings (1-5), **99 % ratings missing**
- \$1 M prize for 10 % reduction in RMSE over Cinematch (a method developed by Netflix)
- BellKor's Pragmatic Chaos declared winners on 9/21/2009
- used ensemble of models including **low-rank factorization**

	The Game	Revolutionary Road	Tinker Tailor Soldier Spy	The Imitation Game	...
User 1	★★★★★	★★★	?	★★★★★	...
User 2	?	?	★★★★★	★★★★★	...
User 3	★★★	★★★★★	?	★★★★★	...
User 4	★★★★★	?	★★★★★	?	...
...

Motivation: Multi-Sensors Monitoring

- Hundreds of sensors are used for condition monitoring of industrial assets.
- Different **sampling frequencies** of sensors and/or their malfunctions will result in missing data.
- Goal is **impute** the missing values based on the temporal and spatial correlation of sensors.
- This is achievable if the matrix is low rank



	Sensor 1	Sensor 2	Sensor 3	Sensor 4	...
Engine 1					...
Engine 2					...
Engine 3					...
...

Matrix Completion Formulation

Assuming the that the matrix is low rank, the missing values can be imputed by solving:

$$\begin{aligned} & \min\{\text{rank}(Z)\} \\ \text{subject to: } & Z_{ij} = X_{ij}; (i,j) \in \Omega_{\text{observed}} \end{aligned}$$

- Use the following projection function $P_{\Omega}(X)_{n \times m}$:

$$P_{\Omega}(X)_{ij} = \begin{cases} X_{ij} & \text{if } (i,j) \text{ is observed} \\ 0 & \text{if } (i,j) \text{ is missing} \end{cases}$$



$$\begin{aligned} & \min\{\text{rank}(Z)\} \\ \text{subject to: } & P_{\Omega}(Z) = P_{\Omega}(X) \end{aligned}$$

×	?	?	?	×	?
?	?	×	×	?	?
×	?	?	×	?	?
?	?	×	?	?	×
×	?	?	?	?	?
?	?	×	×	?	?

Matrix Completion Formulation

Assuming the that the matrix is low rank, the missing values can be imputed by solving:

$$\begin{aligned} & \min\{\text{rank}(Z)\} \\ \text{subject to: } & P_{\Omega}(Z) = P_{\Omega}(X) \end{aligned}$$

Convex Relaxation



$$\begin{aligned} & \min\{\|Z\|_*\} \\ \text{subject to: } & P_{\Omega}(Z) = P_{\Omega}(X) \end{aligned}$$

×	?	?	?	×	?
?	?	×	×	?	?
×	?	?	×	?	?
?	?	×	?	?	×
×	?	?	?	?	?
?	?	×	×	?	?

Singular Value Thresholding (SVT)

Consider the **singular value decomposition (SVD)** of a matrix $X \in \mathbb{R}^{n_1 \times n_2}$ of rank r .

$$\begin{aligned} X &= U \Sigma V^T, \\ \Sigma &= \text{diag}(\{\sigma_i\}_{1 \leq i \leq r}), \end{aligned}$$

where $U \in \mathbb{R}^{n_1 \times r}$ and $V \in \mathbb{R}^{n_2 \times r}$ matrices with orthogonal columns and σ_i are positive singular values.

For each $\tau \geq 0$, the **soft-thresholding operator, S_τ** is defined by

$$\begin{aligned} S_\tau(X) &= U S_\tau(\Sigma) V^T, \\ S_\tau(\Sigma) &= \text{diag}\{(\sigma_i - \tau)_+\}, \end{aligned}$$

where $(a)_+ = \max(0, a)$.

Matrix Completion using SVT

SVT can be used to solve the relaxed convex optimization problem:

$$\begin{aligned} &\min\{\|Z\|_*\} \\ &\text{subject to: } P_\Omega(Z) = P_\Omega(X) \end{aligned}$$

For a fixed $\tau > 0$ and a sequence of $\{\delta_k\}$ of positive step sizes, start with $Y_0 = 0 \in \mathbb{R}^{n_1 \times n_2}$, and for $k = 1, 2, \dots$ iteratively compute

$$\begin{cases} Z^k = S_\tau(Y^{k-1}) \\ Y^k = Y^{k-1} + \delta_k P_\Omega(X - Z^k) \end{cases}$$

Until convergence (i.e., a stopping criterion is reached).

Matrix Completion with Noisy Data

In real life, however, observations are noisy. Therefore, the equality constraint should be relaxed:

$$\begin{aligned} & \min\{\|Z\|_*\} \\ \text{subject to: } & \|P_\Omega(Z) - P_\Omega(X)\|_F^2 < \epsilon \end{aligned}$$

The Langrangian equivalent of the problem is

$$\min\left\{\lambda\|Z\|_* + \frac{1}{2}\|P_\Omega(Z) - P_\Omega(X)\|_F^2\right\},$$

For some $\lambda > 0$.

$$\begin{bmatrix} \times & ? & ? & ? & \times & ? \\ ? & ? & \times & \times & ? & ? \\ \times & ? & ? & \times & ? & ? \\ ? & ? & \times & ? & ? & \times \\ \times & ? & ? & ? & ? & ? \\ ? & ? & \times & \times & ? & ? \end{bmatrix}$$

Matrix Completion with Noisy Data

$$\min\left\{\lambda\|Z\|_* + \frac{1}{2}\|P_\Omega(Z) - P_\Omega(X)\|_F^2\right\},$$

If the PFBS (Proximal Forward-Backward Splitting) method is used, the following iterative algorithm for solving this problem can be obtained:

$$\begin{cases} Z^k = S_{\lambda\delta_{k-1}}(Y^{k-1}), \\ Y^k = Z^k + \delta_k P_\Omega(X - Z^k) \end{cases}$$

Until convergence (i.e., a stopping criterion is reached).

Some theoretical results can be found in Ref. 1 and 2.

Example

```
%Generating original matrix
n1 = 10; n2 = 8;
A = randi([-20,20],n1,n2);
r = 2;
[U, S, V] = svd(A);
if n1 < n2
    s = diag(S); s(r+1:end)=0; S=[diag(s) zeros(n1,n2-n1)];
else
    s = diag(S); s(r+1:end)=0; S=[diag(s); zeros(n1-n2,n2)];
end

X = U* S* V';
X0 = X;

%Removing 20% of the observations
A = [rand(n1,n2)>0.80];
X(A) = 0;
m = sum(sum(A==0));
```

$X =$


2.466	1.153	-7.324	1.214	-1.517	-0.304	-2.04	2.383
11.09	-7.026	0.885	-5.155	-3.33	14.58	11.08	4.589
2.521	6.269	-21.59	5.667	-3.008	-6.959	-10.53	4.991
-6.835	6.415	-6.321	4.99	1.456	-11.71	-10.29	-1.782
-12.75	2.116	15.49	0.744	5.532	-8.97	-2.851	-8.264
4.237	-4.477	5.305	-3.528	-0.759	7.909	7.21	0.853
16.03	-5.505	-11.6	-3.408	-6.143	14.99	8.305	8.966
7.485	-7.445	8.087	-5.829	-1.474	13.37	11.97	1.741
-8.605	-3.455	23.99	-3.744	5.132	0.321	6.179	-8.03
-0.696	4.11	-10.22	3.513	-0.841	-5.706	-6.784	1.553

2.466	1.153	-7.324	1.214	-1.517	-0.304	-2.04	2.383
11.09	-7.026	0.885	-5.155	-3.33	14.58	11.08	4.589
0	6.269	-21.59	5.667	-3.008	-6.959	0	4.991
-6.835	6.415	0	4.99	1.456	-11.71	-10.29	-1.782
-12.75	2.116	15.49	0.744	5.532	-8.97	-2.851	0
4.237	-4.477	5.305	-3.528	-0.759	7.909	0	0.853
16.03	-5.505	-11.6	-3.408	0	14.99	8.305	8.966
0	-7.445	8.087	-5.829	-1.474	13.37	0	0
0	-3.455	23.99	0	5.132	0	6.179	0
-0.696	4.11	-10.22	3.513	0	-5.706	-6.784	1.553



```
%Initialization
Y=zeros(n1,n2);
delta = n1*n2/m;
tau = 250;

%Iterations
vec = zeros(500,1);
for i = 1:500
    [U, S, V] = svd(Y);
    S_t = (S-tau);
    S_t(S_t<0) = 0;
    Z = U*S_t*V';
    P = X-Z;
    P(A) = 0;
    Y0 = Y;
    Y = Y0 + delta*P;
    vec(i) = sum(sum((Y-Y0).^2));
    err(i) = sum(sum((X0-Z).^2))/sum(sum((X0).^2));
end
```


```
% plot the results
figure;plot(vec);
figure;plot(err);
figure;
Ar=reshape(A, n1*n2,1);
Xr=reshape(X0, n1*n2,1);Xr=Xr(Ar);
Zr=reshape(Z, n1*n2,1);Zr=Zr(Ar);
subplot(2,1,1);plot(Xr);hold on;plot(Zr, 'r');
subplot(2,1,2);plot(Xr-Zr);
```



$$\begin{cases} Z^k = S_\tau(Y^{k-1}) \\ Y^k = Y^{k-1} + \delta_k P_\Omega(X - Z^k) \end{cases}$$



$$\|X - Z\|_2^2 = 4.6261e - 09$$



References

- [1] Jian-feng Cai, Emmanuel J. Candes, And Zuowei Shen. A Singular Value Thresholding Algorithm For Matrix Completion. Technical Report.
- [2] Emmanuel J. Candès and Benjamin Recht. Exact Matrix Completion via Convex Optimization.
- [3] Trevor Hastie, Rahul Mazumder, Jason D. Lee, Reza Zadeh, Matrix Completion and Low-Rank SVD via Fast Alternating Least Squares.

Topics on High-Dimensional Data Analytics

Regularization Applications

Kamran Paynabar, Ph.D.

Associate Professor
School of Industrial and Systems Engineering

Robust PCA



Learning Objectives

- Explain difference between PCA and Robust PCA.
- Use nuclear regularization to formulate the Robust PCA problem.
- Apply optimization algorithms to solve the Robust PCA problem.




Traditional PCA

- Principal Component Analysis reduces data dimensions by finding a **low-rank** representation of data.
- Suppose $X = [x_1 \ x_2 \ \dots \ x_p]$ is the centered matrix data.
- A low-rank representation of X is given by $X = L + E$; where L is a low-rank matrix and E is noise.
- The low-rank subspace L can be found by solving

$$\arg \min_L \|X - L\|_F^2 \quad s.t. \text{rank}(L) \leq k$$


- This problem can be solved via **truncated SVD**:

$$L = U\Sigma V' = \sum_{i \leq k} \sigma_i \mathbf{u}_i \mathbf{v}_i'$$



PCA in Presence of Outliers


- Similar to regression, PCA is highly sensitive to **outliers**.



- Assuming outliers are sparse, **Robust PCA** can retrieve the correct low-rank structure

Robust PCA Applications

- Face recognition
- Anomaly detection
- Denoising
- Text mining
- Web mining
- Image and video repair
- Video Surveillance



Robust PCA

$$M = L + S$$

- M is the **observed** data matrix
- L is a **low-rank** matrix (to be estimated)
- S is a matrix of **sparse** outliers (to be estimated)

By finding L and S , we can retrieve **low-dimensional linear structure** from non-ideal observations.

Robust PCA Formulation

$$M = L + S$$

- L is **low-rank**
- S is **sparse**
- L and S could be found by solving

$$\min\{\text{rank}(L) + \lambda \|S\|_0\}$$

$$\text{subject to: } M = L + S$$

$$\text{rank}(L) = \#\{\sigma(L) \neq 0\} \quad \|S\|_0 = \#\{S_{ij} \neq 0\}$$

Intractable!!!!

Robust PCA Formulation

$$\min\{\text{rank}(L) + \lambda\|S\|_0\}$$

$$\text{subject to: } M = L + S$$

$$\text{rank}(L) = \#\{\sigma(L) \neq 0\}$$

$$\|S\|_0 = \#\{S_{ij} \neq 0\}$$



Convex Relaxation



$$\|L\|_* = \sum_i \sigma(L)$$

Nuclear norm: sum of singular values

$$\|S\|_1 = \sum_{i,j} |S_{ij}|$$

L₁ norm: sum of absolute values

This is also known as Principal Component Pursuit (PCP)

See Chandrasekaran, Sanghavi, Parrilo, Willsky ('09)

Solving RPCA

- Review of **Augmented Lagrangian Multipliers** Method:

$$\min f(X), \quad \text{subject to } h(X) = 0,$$

$$L(X, Y, \mu) = f(X) + \langle Y, h(X) \rangle + \frac{\mu}{2} \|h(X)\|_F^2$$

(General Method of Augmented Lagrange Multiplier)

- 1: $\mu \geq 1$.
 - 2: while not converged do
 - 3: Solve $X_{k+1} = \underset{X}{\operatorname{argmin}} L(X, Y_k, \mu_k)$.
 - 4: $Y_{k+1} = Y_k + \mu_k h(X_{k+1})$;
 - 5: Update μ_k to μ_{k+1} .
 - 6: end while
- Output: X_k .
-

Solving RPCA

- Augmented Lagrangian Multiplier form:

$$l(L, S, Y) = \|L\|_* + \lambda \|S\|_1 + \langle Y, M - L - S \rangle + \frac{\mu}{2} \|M - L - S\|_F^2.$$

$$l(L, S, Y; \mu) = \|L\|_* + \lambda \|S\|_1 + \frac{\mu}{2} \left\| M - L - S + \frac{Y}{\mu} \right\|_F^2 + \frac{\mu}{2} \left\| \frac{Y}{\mu} \right\|_F^2$$

Main Idea:

- Given S and Y, Update L

$$\arg \min_L \|L\|_* + \frac{\mu}{2} \left\| M - L - S + \frac{Y}{\mu} \right\|_F^2 \quad \xrightarrow{X = M - S + \frac{Y}{\mu}} \quad L = D_{1/\mu}(X) = US_{1/\mu}(\Sigma)V^T$$

- Given L and Y, Update S

$$\arg \min_S \lambda \|S\|_1 + \frac{\mu}{2} \left\| M - L - S + \frac{Y}{\mu} \right\|_F^2 \quad \xrightarrow{\quad} \quad S_{ij} = S_{\lambda/\mu}(X) = \text{sgn}(X) \max(|X| - \frac{\lambda}{\mu}, 0)$$

- Given L and S, Update Y

$$Y_{k+1} = Y_k + \mu(M - L - S)$$

- Iteration until Convergence

Solving RPCA - Algorithm

$$l(L, S, Y) = \|L\|_* + \lambda \|S\|_1 + \langle Y, M - L - S \rangle + \frac{\mu}{2} \|M - L - S\|_F^2.$$

- Algorithm:

```

1: initialize:  $S_0 = Y_0 = 0, \mu > 0.$ 
2: while not converged do
3:   compute  $L_{k+1} = \mathcal{D}_{1/\mu}^1(M - S_k - \mu^{-1}Y_k);$            (soft-thresholding on singular values by  $\frac{1}{\mu}$ )
4:   compute  $S_{k+1} = \mathcal{S}_{\lambda/\mu}^\mu(M - L_{k+1} + \mu^{-1}Y_k);$        (soft-thresholding scalar entries by  $\frac{\lambda}{\mu}$ )
5:   compute  $Y_{k+1} = Y_k + \mu(M - L_{k+1} - S_{k+1});$ 
6: end while
7: output:  $L, S.$ 


```

Example

Input




Output



```
D = double(rgb2gray(imread('building.png')));
lambda=1e-2;
[m n] = size(D); tol = 1e-7; maxIter = 1000;
% Initialize A,E,Y,u
Y = D; norm_two = norm(Y); norm_inf = norm(Y(:), inf) / lambda;
Y = Y / norm_inf;
A_hat = zeros(m, n); E_hat = zeros(m, n);
mu = 1.25/norm_two; mu_bar = mu * 1e7; rho = 1.5; % this one can be tuned
d_norm = norm(D, 'fro');
iter = 0; total_svd = 0; converged = false; stopCriterion = 1;
while ~converged
    iter = iter + 1;
    temp_T = D - A_hat + (1/mu)*Y;
    E_hat = max(temp_T - lambda/mu, 0);
    E_hat = E_hat + min(temp_T + lambda/mu, 0);
    [U S V] = svd(D - E_hat + (1/mu)*Y, 'econ');
    diagS = diag(S);
    svp = length(find(diagS > 1/mu));
    A_hat = U(:, 1:svp) * diag(diagS(1:svp) - 1/mu) * V(:, 1:svp);
    total_svd = total_svd + 1;
    Z = D - A_hat - E_hat;
    Y = Y + mu*Z;
    mu = min(mu*rho, mu_bar);
    % stop Criterion
    stopCriterion = norm(Z, 'fro') / d_norm;
    if stopCriterion < tol
        converged = true;
    end
end
```

Liang, Ren, Zhang, and Ma, ECCV 2012

Robust PCA for Noisy Data



$$M = L + S + E$$

- L is **low-rank**, S is **sparse**, E is **noise**
- L and S could be found by solving

$$\min\{\|L\|_* + \lambda\|S\|_1\}$$

$$\text{subject to: } \|M - (L + S)\|_F < \delta$$

References

- [1] Martin J Wainwright. Structured regularizers for high-dimensional problems: Statistical and computational issues. *Annual Review of Statistics and Its Application*, 1:233–253, 2014.
- [2] Sahand Negahban, Bin Yu, Martin J Wainwright, and Pradeep K Ravikumar. A unified framework for high-dimensional analysis of m-estimators with decomposable regularizers. In *Advances in Neural Information Processing Systems*, pages 1348–1356, 2009.
- [3] Sparse Representation and Low-Rank Representation for Biometrics -- Theory, Algorithms, and Applications, ICB 2013 Short Course
- [4] John Wright, Arvind Ganesh, Shankar Rao, Yigang Peng, and Yi Ma. Robust principal component analysis: Exact recovery of corrupted low-rank matrices via convex optimization. In *Advances in neural information processing systems*, pages 2080–2088, 2009.
- [5] Emmanuel J Candès, Xiaodong Li, Yi Ma, and John Wright. Robust principal component analysis? *Journal of the ACM (JACM)*, 58(3):11, 2011.

Topics on High-Dimensional Data Analytics

Regularization Applications

Kamran Paynabar, Ph.D.


Associate Professor
School of Industrial and Systems Engineering

Sparse Smooth Decomposition






Learning Objectives

- Define Smooth Sparse Decomposition (SSD)
- Use L1 and L2 regularization to formulate the SSD problem.
- Apply optimization algorithms to solve a the SSD problem.




Motivation

- In practice, high-dimensional data
 - are noisy;
 - often have functional structure that makes them smooth
 - sometimes contain sparse features.
- Decomposition of an observation into these three elements have various applications, including:
 - Denoising
 - Sparse feature learning
 - Image background retrieval
 - Function estimation in presence of outliers and noise
 - Anomaly detection



Sparse Smooth Decomposition (SSD)

- Objective:** To decompose an HD noisy observation into functional mean, sparse features, and noise (Yan et al., 2017): $y = \mu + a + e$



Assumptions


- Functional mean is **smooth**.
- Features are **sparse** features and exhibit a **different structure** from the functional mean

Two-step Approach: Smoothing + Detection

Methods	Edge Based	Thresholding Based	Region Based
Criterion	Look for discontinuity	Look for largest	Look for similarity
Output	Edge	Region	Region
Preprocessing	Smoothing	Smoothing	Smoothing, need a seed point to start
Postprocessing	Close the Edge to get boundary	Clean small areas	Clean small areas
Algorithms	Sobel Edge Detection ^[1] Jump Regression ^[2]	Ostu's Global method ^[3] Nick's Local method ^[4]	Region Growing algorithm ^[5]

Some issues:

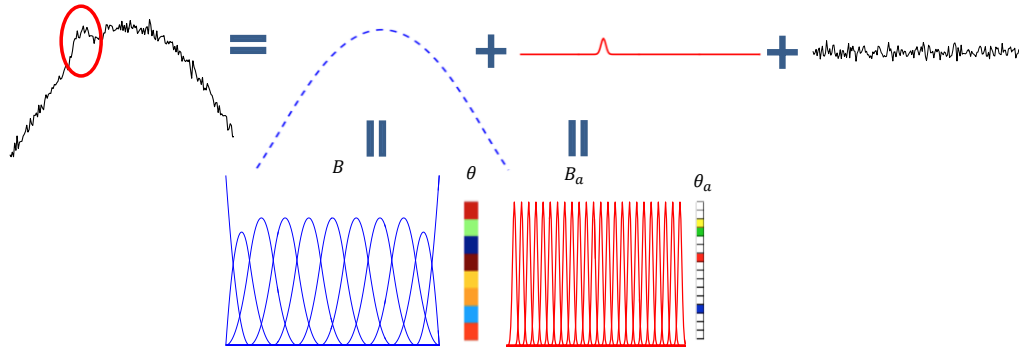
- Smoothing blurs the boundaries of defect regions;
- Smoothing is inaccurate in presence of anomaly;
- Often, a post-processing task is required to connect detected boundaries;



SSD Formulation

Let's begin with 1D case: $y = \mu + a + e \quad \rightarrow \quad y = B\theta + B_a\theta_a + e$

$$\operatorname{argmin}_{\theta, \theta_a} \lambda \theta' R \theta + \gamma \|\theta_a\|_1 + \|e\|^2, s.t. y = B\theta + B_a\theta_a + e$$



R is the roughness matrix that can be defined by $R = D^T D$ where D is the first difference matrix

SSD via Iterative Thresholding Algorithm

$$\operatorname{argmin}_{\theta, \theta_a} \lambda \theta' R \theta + \gamma \|\theta_a\|_1 + \|e\|^2, s.t. y = B\theta + B_a\theta_a + e$$

- Propose an optimization algorithm based on ITA (Daubechies, et al, 2004)

Iterative Thresholding Algorithm

In the k^{th} iteration, update $\mu^{(k)}$ and $\theta_a^{(k)}$ by

$$\mu^{(k)} = H(y - B_a\theta_a^{(k-1)}), H = B(B'B + \lambda I)^{-1}B' \text{ is the projection matrix}$$

$$\theta_a^{(k)} = T_{\tau^{(k)}}(\theta_a^{(k-1)} - c^{(k)}B_a'(B_a\theta_a^{(k-1)} + \mu^{(k)} - y))$$


$$c^{(k)} = \frac{2}{L'} \cdot \tau^{(k)} = \frac{\gamma}{L}$$

- $p = 1$, convex optimization
 - $T(\cdot)$ is the soft-thresholding operator.

Accelerated Algorithm for SSD

Nesterov's acceleration (Beck and Teboulle, 2009)

- $x_a^{(k)} = T_{\tau^{(k)}}(\theta_a^{(k-1)}, c^{(k)}, B_a, \mu^{(k)}, y)$
- $t_{k+1} = (1 + \sqrt{1 + 4t_k^2})/2$
- $\theta_a^{(k+1)} = x_a^{(k)} + \frac{t_k - 1}{t_{k+1}}(x_a^{(k)} - x_a^{(k-1)})$



$$\begin{cases} y_B^{(k)} = H(y - B_S \theta_S^{(k-1)}) \\ \theta_S^{(k)} = \text{softthreshold}(x^{(k-1)} + \frac{2}{L} B_S' (y - B_S \theta_S^{(k-1)} - y_B^{(k)}), \frac{\gamma}{L}) \end{cases}$$

Convergence Rate & Error Bound of ITA Algorithm for SSD

Restricted isometry property

- δ_{3s+1} is RIP constant which satisfies (for all s -sparse vector θ_{a1}, θ_{a2})
 $(1 - \delta_{3s+1}) \|\theta_{a1} - \theta_{a2}\|_2^2 \leq \|(I - H)(B_a \theta_{a1} - B_a \theta_{a2})\|_2^2 \leq (1 + \delta_{3s+1}) \|\theta_{a1} - \theta_{a2}\|_2^2$

Theorem

If $\rho_s = L \delta_{3s+1} ((I - H) B_a) < 1$, the anomaly estimation in the k^{th} iteration of the algorithm with $\lambda = 0$ is bounded by


$$\|\theta_a^{(k)} - \theta_a^*\|_2 \leq (\rho_s)^k \|\theta_a^{(0)} - \theta_a^*\|_2 + \frac{L}{1 - \rho_s} \|B_a^T (I - H) \epsilon\|_2$$

Linear convergence rate


Estimation error

- Noiseless case: $\|\epsilon\|_2 = 0$, perfect recovery
- Noisy case: detect anomaly with $B_a \perp \epsilon$

Generalization of SSD to 2D Case (Images)



$$\underset{\theta, \theta_S}{\operatorname{argmin}} \|\tilde{e}\|^2 + \theta' R \theta + \gamma |\theta_S|_1, \text{ s.t. } y = (B_1 \otimes B_2) \theta + (B_{S,1} \otimes B_{S,2}) \theta_S + \tilde{e}$$



Generalization of SSD to 2D Case (Images)

$$\underset{\theta, \theta_S}{\operatorname{argmin}} \|\tilde{e}\|^2 + \theta' R \theta + \gamma |\theta_S|_1, \text{ s.t. } y = (B_1 \otimes B_2) \theta + (B_{S,1} \otimes B_{S,2}) \theta_S + \tilde{e}$$

$$y: n_1 n_2 \times 1$$

$$H: k_1 k_2 \times k_1 k_2$$

$$\begin{cases} y_B^{(k)} = H(y - B_S \theta_S^{(k-1)}) \\ \theta_S^{(k)} = \text{softthreshold}(x^{(k-1)} + \frac{2}{L} B'_S (y - B_S \theta_S^{(k-1)} - y_B^{(k)}), \frac{\gamma}{L}) \end{cases}$$

APG algorithm

Complexity $O(k_1^3 k_2^3 + 6n_1^2 n_2^2 k_1 k_2)$

$$\text{Define } R = \lambda_1 B'_2 B_2 \otimes D'_1 D_1 + \lambda_2 D'_2 D_2 \otimes B'_1 B_1 + \lambda_1 \lambda_2 D'_2 D_2 \otimes D'_1 D_1$$

$$H = H_2 \otimes H_1 \quad H_i = B_i (B'_i B_i + \lambda_i D'_i D_i)^{-1} B'_i, \quad i = 1, 2 \text{ size } n_i \times n_i$$




APG algorithm

$$\begin{cases} Y_B^{(k)} = H_1(Y - B_{S_1} \Theta_S^{(k-1)} B'_{S_2}) H_2 \\ \Theta_S^{(k)} = \text{softthreshold}(X^{(k-1)} + \frac{2}{L} B'_{S_1} (Y - B_{S_1} X^{(k-1)} B'_{S_2} - Y_B^{(k)}) B_{S_2}, \frac{\gamma}{L}) \end{cases}$$

Complexity $O(k_1^3 + k_2^3 + 6n_1^2 k_1 + 6n_2^2 k_2)$

Example


Simulated Image (Y)




Background MSE: 1×10^{-5}
Anomaly FPR: 1%, FNR: 1%

B_a : Quartic kernel


True Anomaly




Estimated Anomaly



True Background



Estimated Background



Matlab file: SSD.m

References

- [1] Yan, H., Paynabar, K., Shi, J., (2017) Anomaly Detection in Images With Smooth Background via Smooth-Sparse Decomposition, *Technometrics*, Vol 59, Issue 1.

Topics on High-Dimensional Data Analytics

Regularization Applications

Kamran Paynabar, Ph.D.

Associate Professor
School of Industrial and Systems Engineering

RKHS Ridge Kernel Regression



Learning Objectives

- Define Reproducing Kernel Hilbert Space
- Use regularization to formulate the ridge regression problem.
- Apply optimization algorithms to solve a the ridge regression problem.



Ridge Regression

Assuming the observations are centered, Ridge estimates can be computed by

$$\arg \min \frac{1}{n} \sum_{i=1}^n (y_i - x_i^T \theta)^2 + \lambda \|\theta\|^2$$

$$\hat{\theta} = (X^T X + \lambda I)^{-1} X^T y$$

- $\lambda \geq 0$ is the tuning parameter controls the amount of shrinkage
- λ is chosen based on some prediction criteria (MSE) using CV or independent validation data set

Ridge Regression

Evaluate ridge regression for a new test point x :

$$\hat{f}(x) = x^T \hat{\theta} = x^T (X^T X + \lambda I)^{-1} X^T y = \underbrace{x^T X^T}_{\text{inner product}} \underbrace{(X^T X + \lambda I)^{-1} y}_{\text{ridge estimate}} = y^T (X X^T + \lambda I)^{-1} X x$$


Only depends on the inner product of predictors

The equality above is true because of the Matrix inversion lemma:

$$X^T (X X^T + \lambda I)^{-1} = (X^T X + \lambda I)^{-1} X^T$$

Kernel as Inner Product

- $k(x, y)$ equivalent to first compute feature $\phi(x)$, and then perform inner product $k(x, y) = \phi(x)^\top \phi(y)$
- A dataset $D = \{x_1, x_2, x_3 \dots x_n\}$
- Compute pairwise kernel function $k(x_i, x_j)$ and form a $n \times n$ kernel matrix (Gram matrix)



$$K = \begin{pmatrix} k(x_1, x_1) & \dots & k(x_1, x_n) \\ \vdots & \ddots & \vdots \\ k(x_n, x_1) & \dots & k(x_n, x_n) \end{pmatrix}$$

Kernel Ridge Regression

$\hat{f}(x) = \hat{\theta}^T x = y^T (XX^T + \lambda I)^{-1} Xx$ only depends on inner products

$$XX^T = \begin{pmatrix} x_1 x_1^T & \dots & x_1 x_n^T \\ \vdots & \ddots & \vdots \\ x_n x_1^T & \dots & x_n x_n^T \end{pmatrix} \quad Xx = \begin{pmatrix} x_1 x \\ \vdots \\ x_n x \end{pmatrix}$$

x_i is the row vector corresponding to the i^{th} row of the matrix X .

Kernel ridge regression: replace the inner products by a kernel function

$$\begin{aligned} XX^T &\rightarrow K = [k(x_i, x_j)]_{n \times n} \\ Xx &\rightarrow k_x = [k(x_i, x)]_{n \times 1} \\ \hat{f}(x) &= \hat{\theta}^T x = y^T (K + \lambda I)^{-1} k_x \end{aligned}$$

Kernel Ridge Regression - Example


```

clc
clear all
Xtrain = (1:100)/100;
Yt = sin(Xtrain*10)+(Xtrain*2).^2;
Ytrain = Yt + 0.2*randn(1,100);
N = 1000;
Xtest = linspace(min(Xtrain),max(Xtrain),N);
Xtrain=Xtrain(:); Ytrain=Ytrain(:); n =
length(Xtrain);Xtest=Xtest(:);

lambda = 0.04;
c = 0.04;

kernel1 = exp(-dist(Xtrain').^2 ./ (2*c));
kernel2 = exp(-pdist2(Xtrain, Xtest).^2 ./ (2*c));
yhatRBF = Ytrain' * ((kernel1 + lambda *
eye(size(kernel1)))) \ kernel2;

```




Another Look to Ridge Regression

- Suppose a set of observation, $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)$ are given. To estimate a regression function, we minimize the following empirical loss function:

$$\arg \min \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2 + \lambda \|f\|_{\mathcal{H}}^2$$

Where \mathcal{H} is the Reproducible Kernel Hilbert Space.

Loosely speaking A **Hilbert space** is a (possibly) infinite dimensional linear space established with a dot product.



Representer Theorem

- The minimizer over the RKHS \mathcal{H} , of the regularized empirical loss function:

$$\arg \min_f \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2 + \lambda \|f\|_{\mathcal{H}}^2$$

Can be represented by the expression

$$f^\lambda(x) = \sum_{i=1}^n \alpha_i K(x_i, x)$$

- Therefore, minimizing over the possibly infinite dimensional Hilbert space boils down to minimizing over \mathbb{R}^n .

RKHS

- Loosely speaking, a **Hilbert space** is a (possibly) infinite dimensional linear space established with a dot product.
- For example, **Square integrable functions** $L_2[a, b]$ (i.e., the integral of the squared function is finite), is Hilbert space and the norm is calculated by

$$\langle f, g \rangle = \int_a^b f(x)g(x)dx$$

- $k(\cdot, \cdot)$ is a **Reproducing Kernel Hilbert Space (RKHS)** of \mathcal{H} if $\forall f \in \mathcal{H}$,
 $f(x) = \langle k(x, \cdot), f(\cdot) \rangle$.

Kernel Ridge Regression

According to the Representer Theorem $\hat{f} = \operatorname{argmin}_{f \in \mathcal{H}} \frac{1}{2} \sum_{i=1}^n (y^{(i)} - f(x^{(i)}))^2 + \frac{\lambda}{2} \|f\|_{\mathcal{H}}^2$

Solved by: $\hat{f}(\cdot) = \sum_{j=1}^n \alpha_j \mathbb{K}(\cdot, x^{(j)})$

$$\begin{aligned} \left\| \sum_{j=1}^n \alpha_j \mathbb{K}(\cdot, x^{(j)}) \right\|_{\mathcal{H}}^2 &= \sum_{i,j=1}^n \alpha_i \alpha_j \langle \mathbb{K}(\cdot, x^{(i)}), \mathbb{K}(\cdot, x^{(j)}) \rangle_{\mathcal{H}} \\ &= \sum_{i,j=1}^n \alpha_i \alpha_j \mathbb{K}(x^{(i)}, x^{(j)}) = \sum_{i,j=1}^n \alpha_i \alpha_j K_{ij} \\ &= \alpha^T \mathbf{K} \alpha \end{aligned}$$

Thus,

$$\hat{\alpha} = \operatorname{argmin}_{\alpha \in \mathbb{R}^n} \frac{1}{2} \|y - \mathbf{K}\alpha\|_2^2 + \frac{\lambda}{2} \alpha^T \mathbf{K} \alpha$$

$$\hat{\alpha} = (\mathbf{K} + \lambda \mathbf{I})^{-1} y$$